



# **Introduction to Kubernetes**


# Features



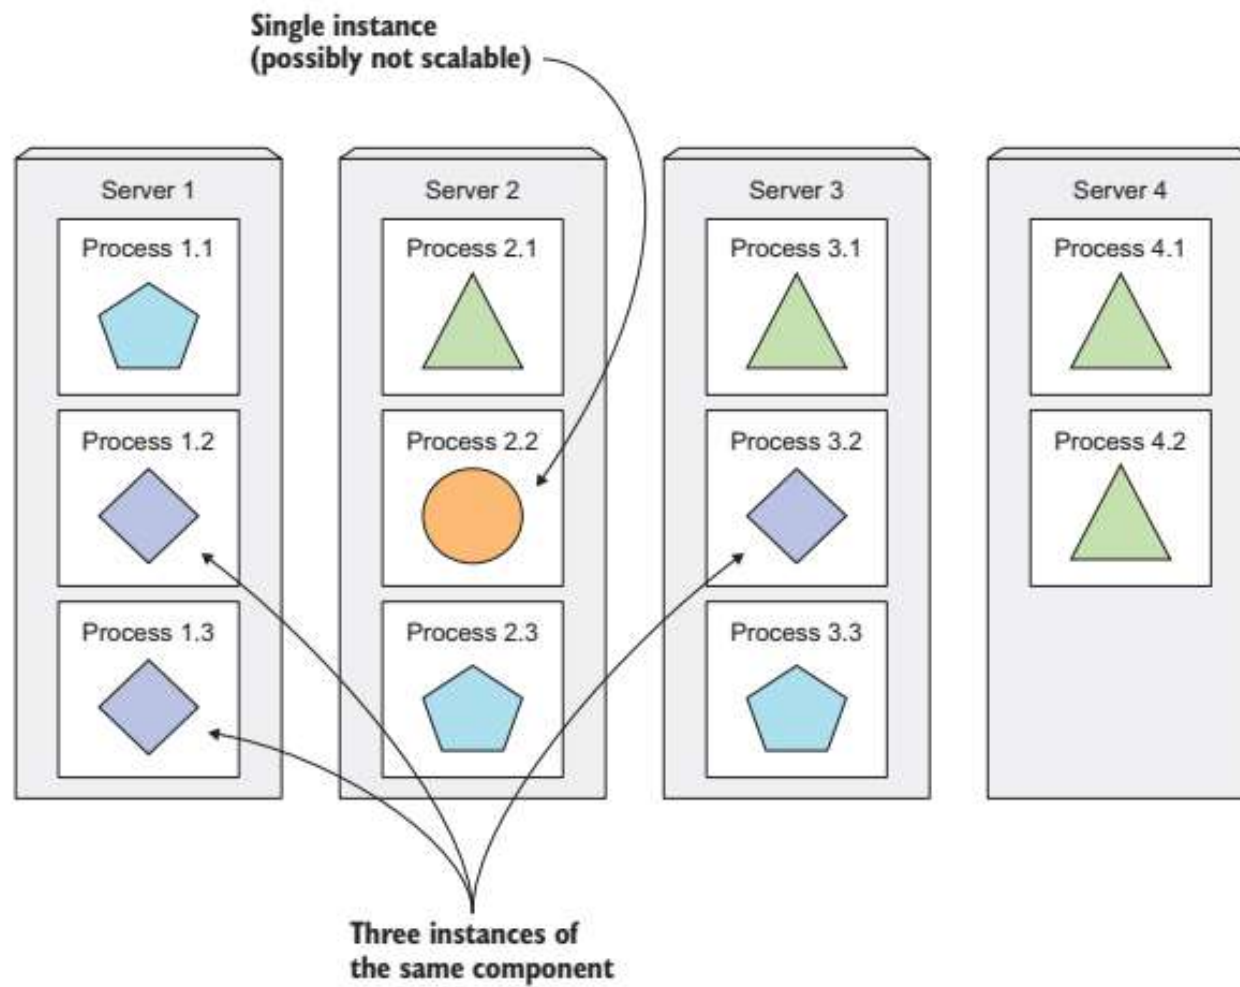
## Monolithic application

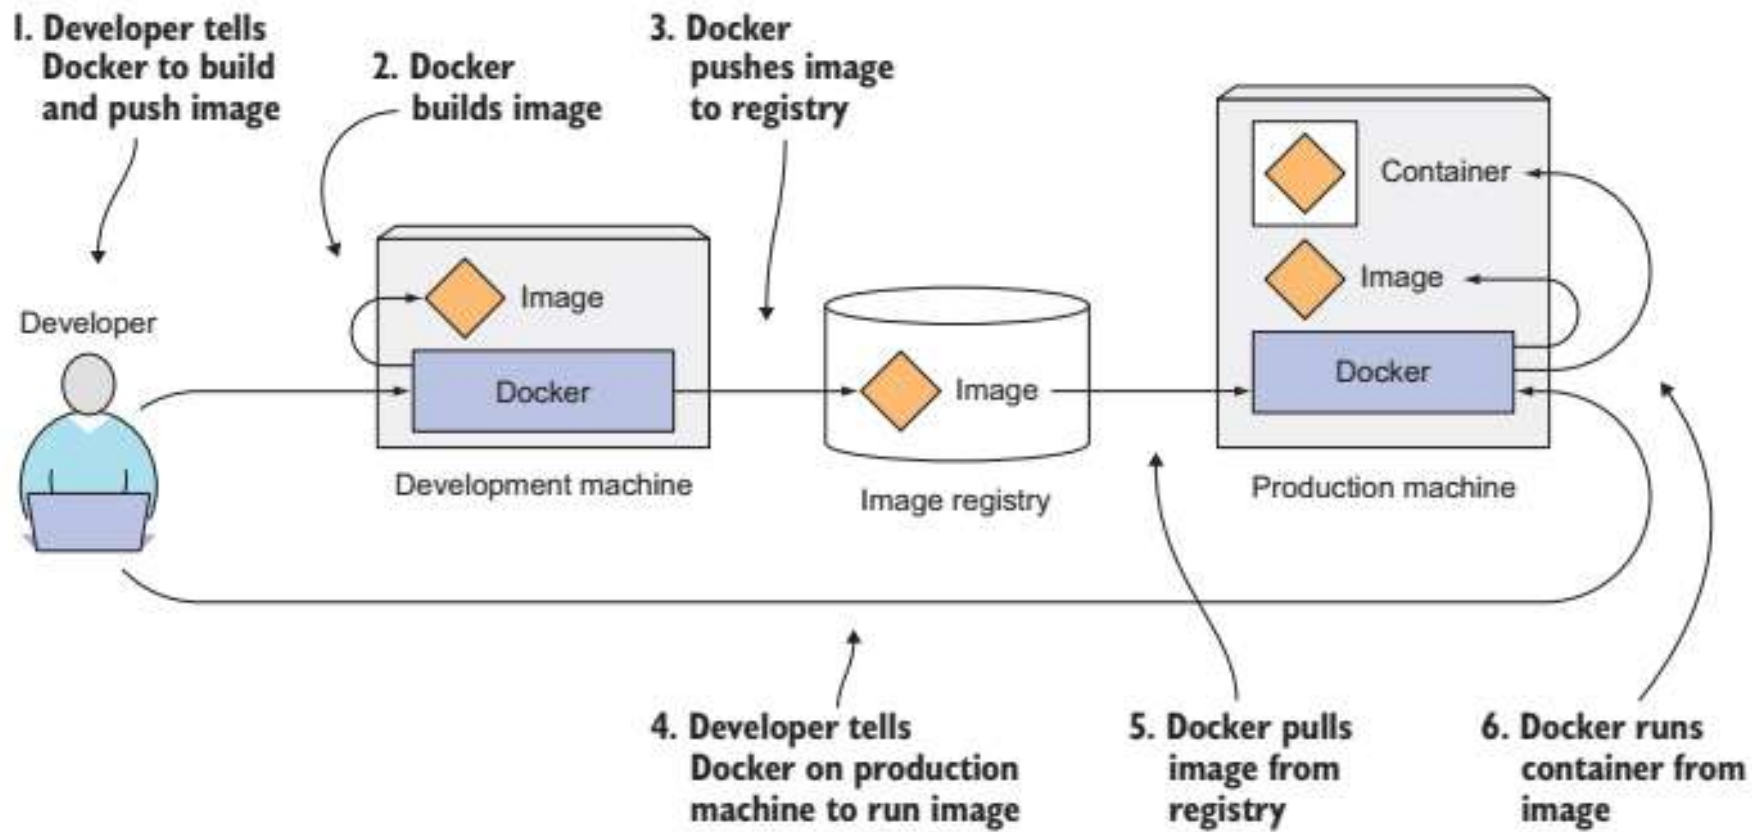


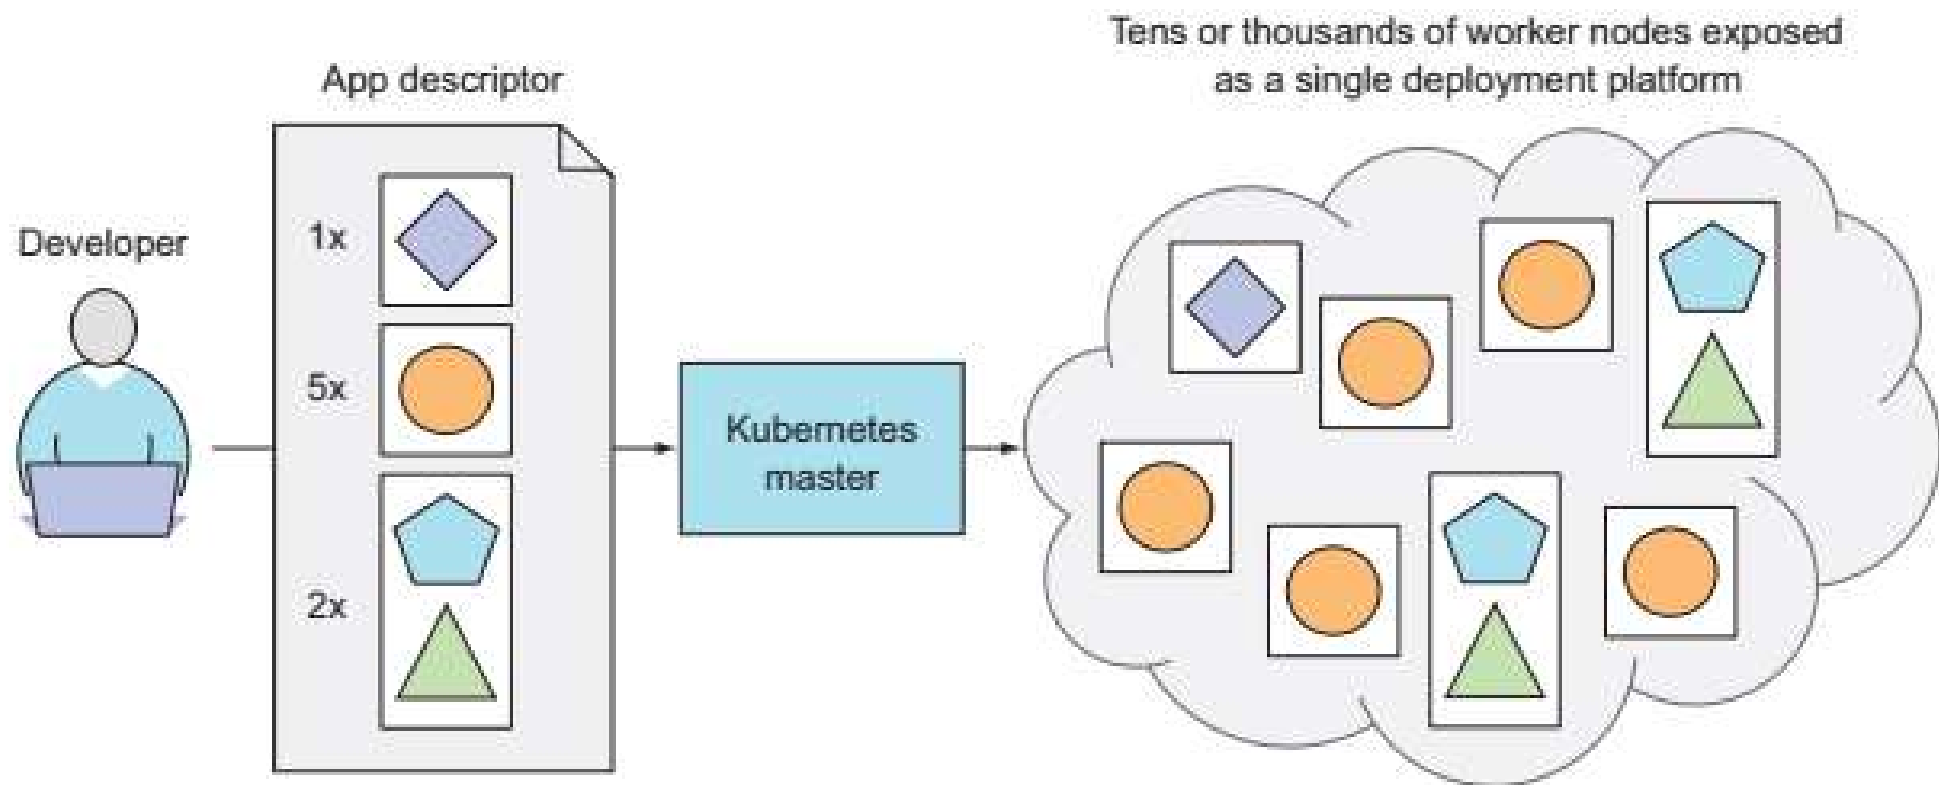
## Microservices-based application




**Figure 1.1** Components inside a monolithic application vs. standalone microservices



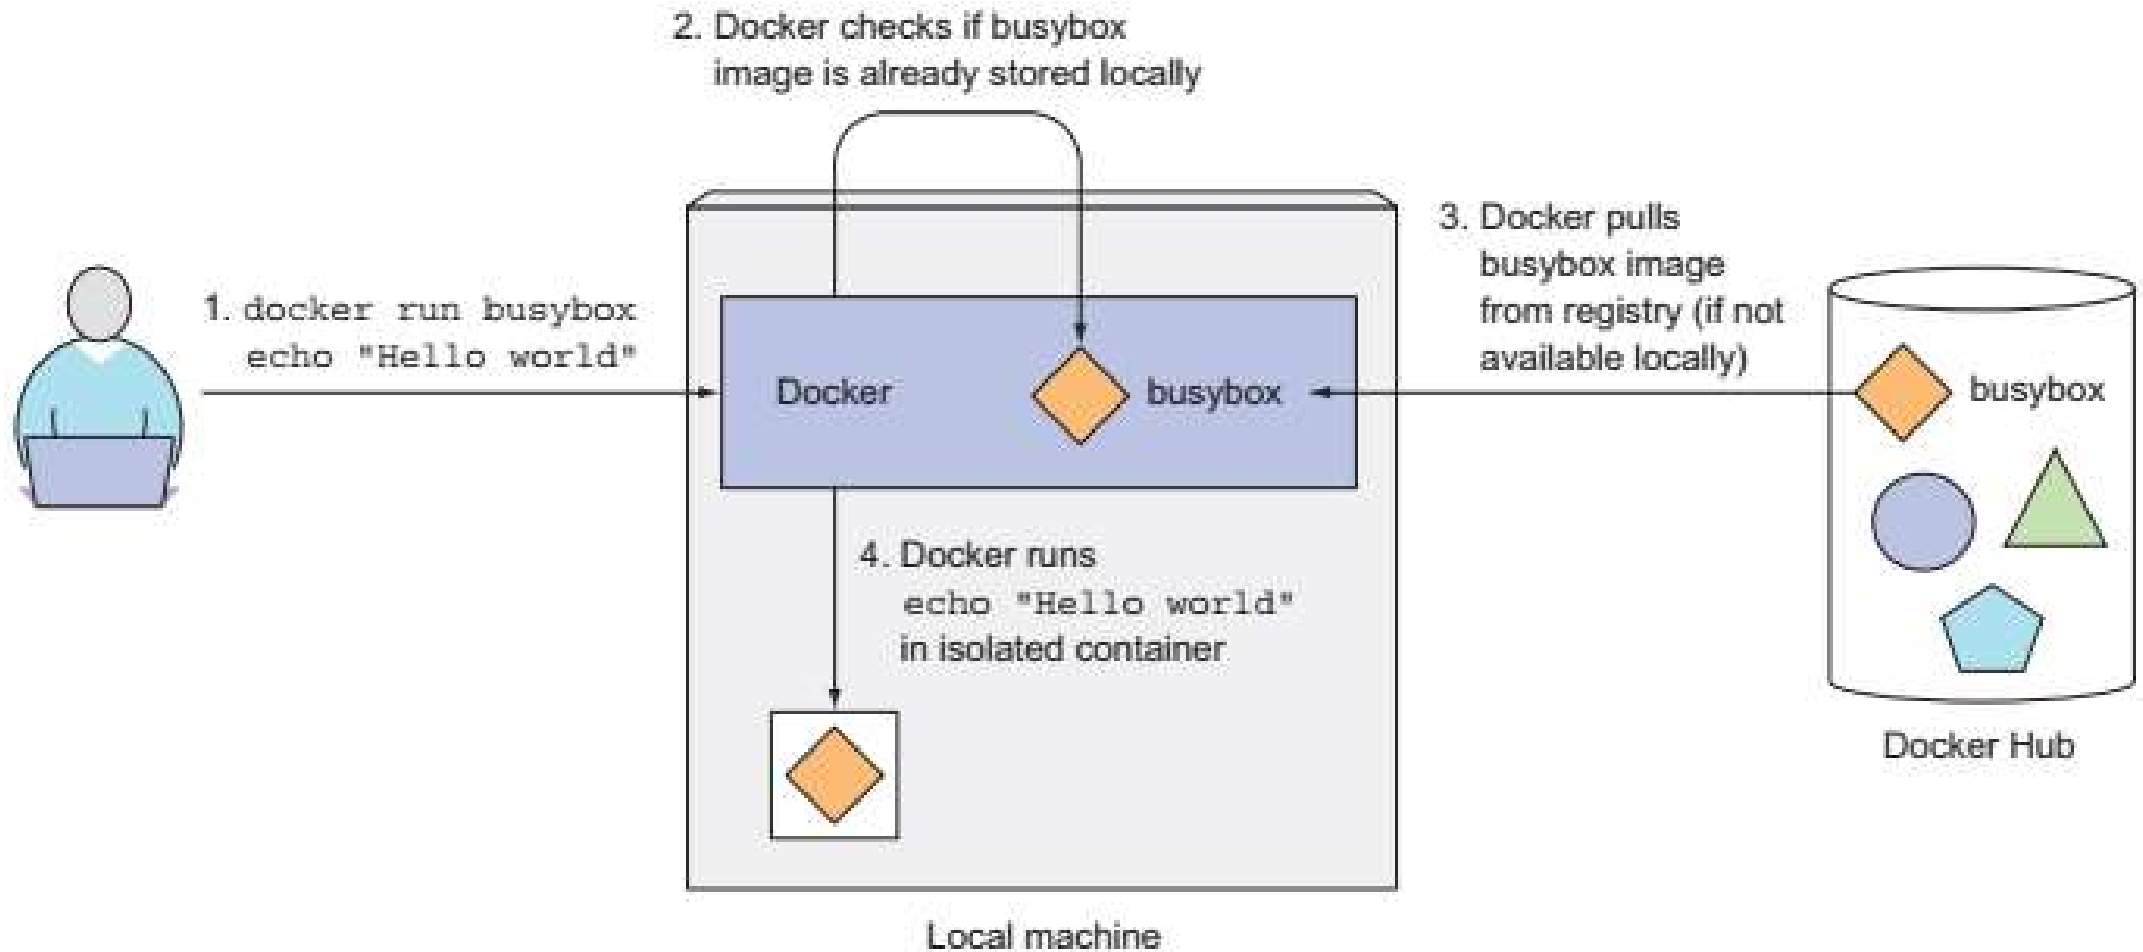




**Figure 1.8** Kubernetes exposes the whole datacenter as a single deployment platform.




**Figure 1.10** A basic overview of the Kubernetes architecture and an application running on top of it

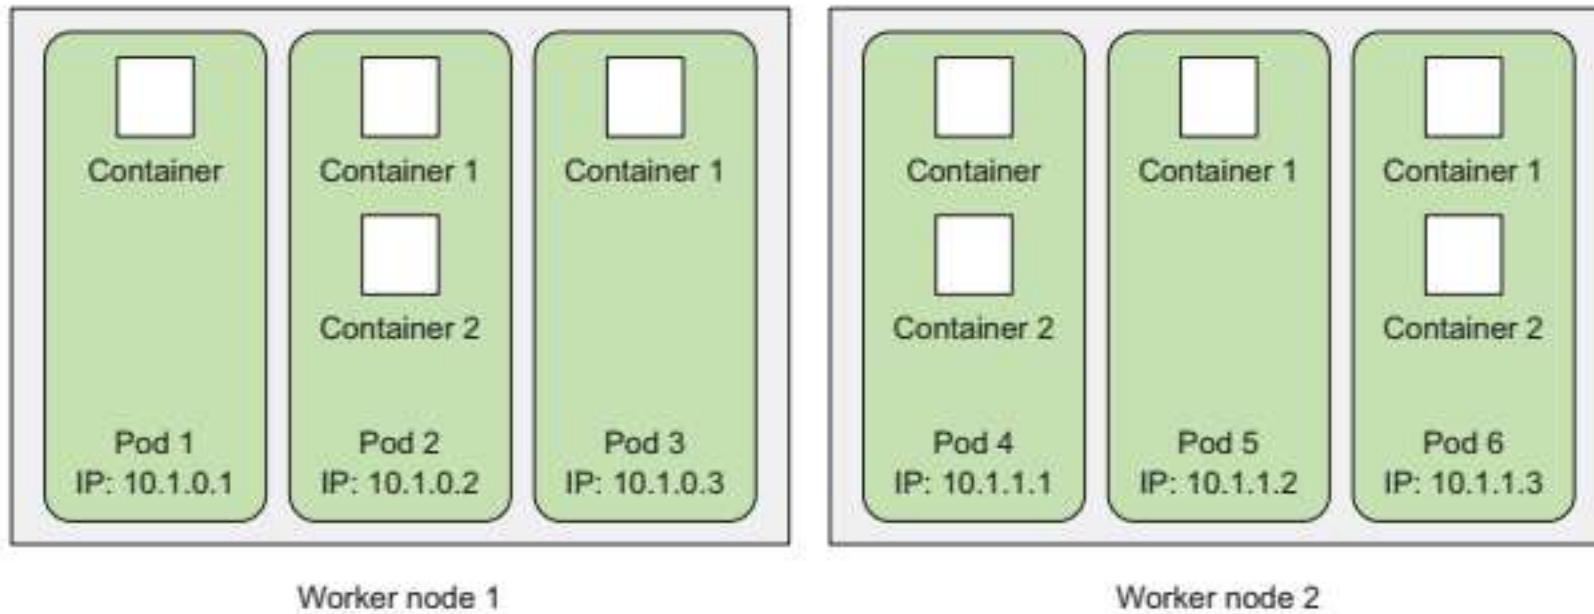


**Figure 2.1** Running echo "Hello world" in a container based on the busybox container image





**Figure 2.4** How you're interacting with your three-node Kubernetes cluster



**Figure 2.5** The relationship between containers, pods, and physical worker nodes


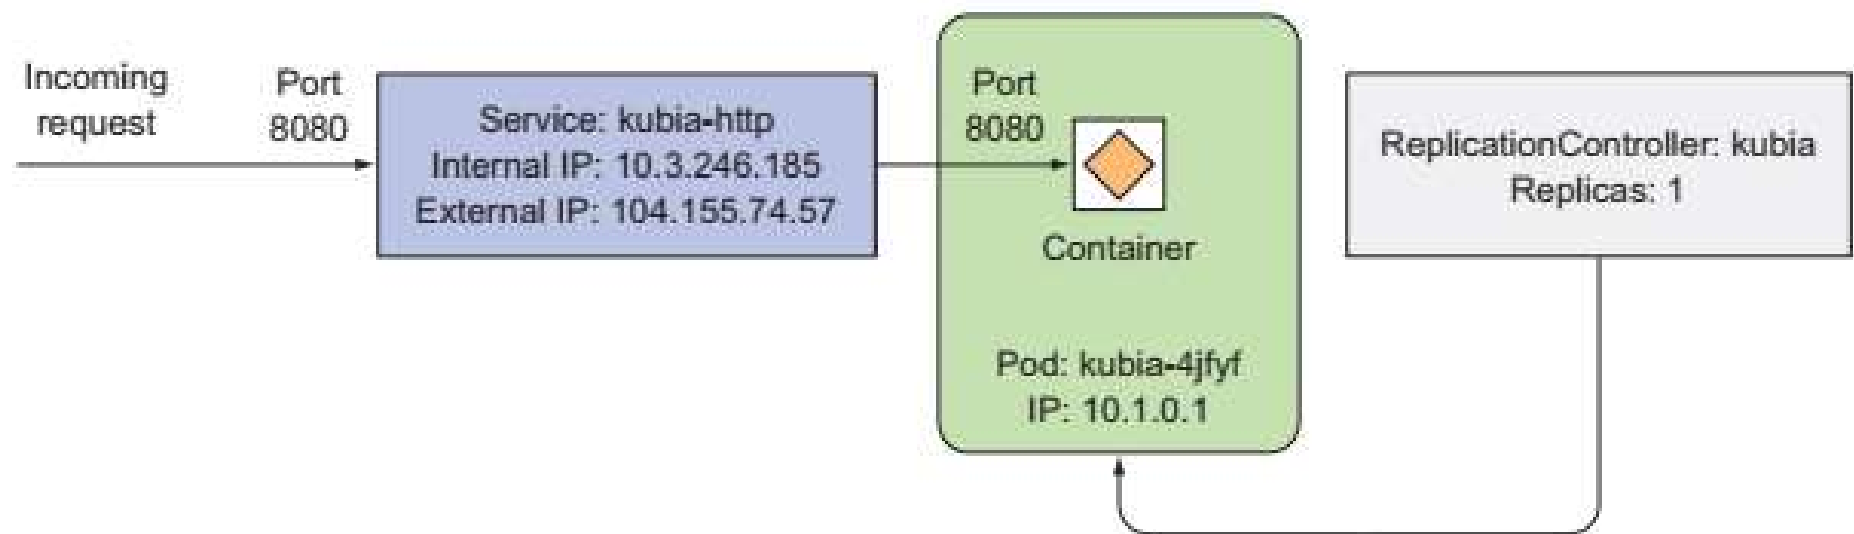
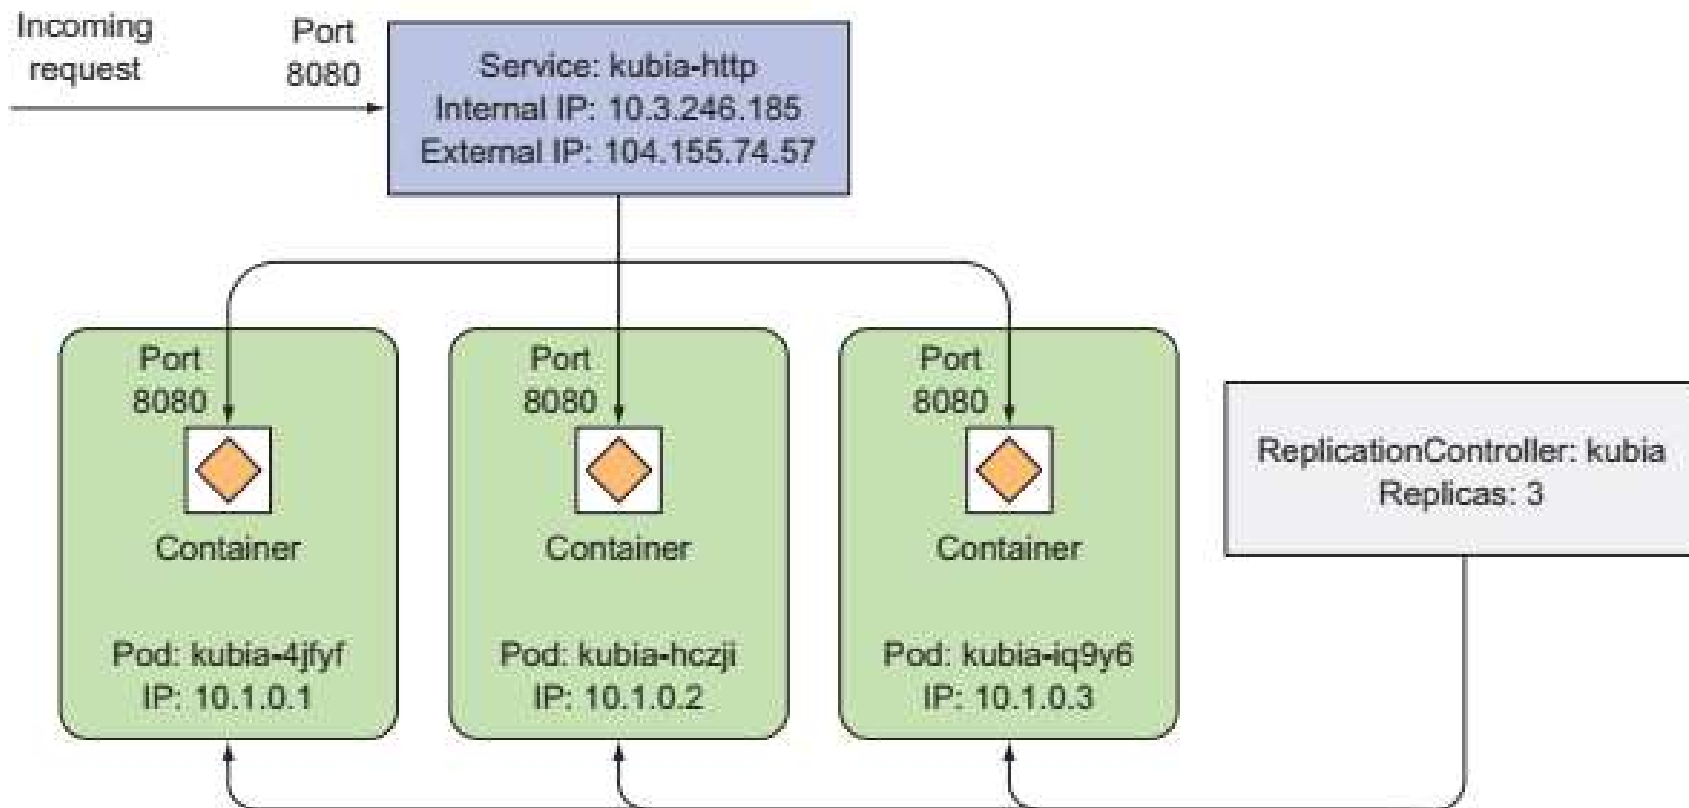



Figure 2.6 Running the luksa/kubia container image in Kubernetes



**Figure 2.7** Your system consists of a ReplicationController, a Pod, and a Service.

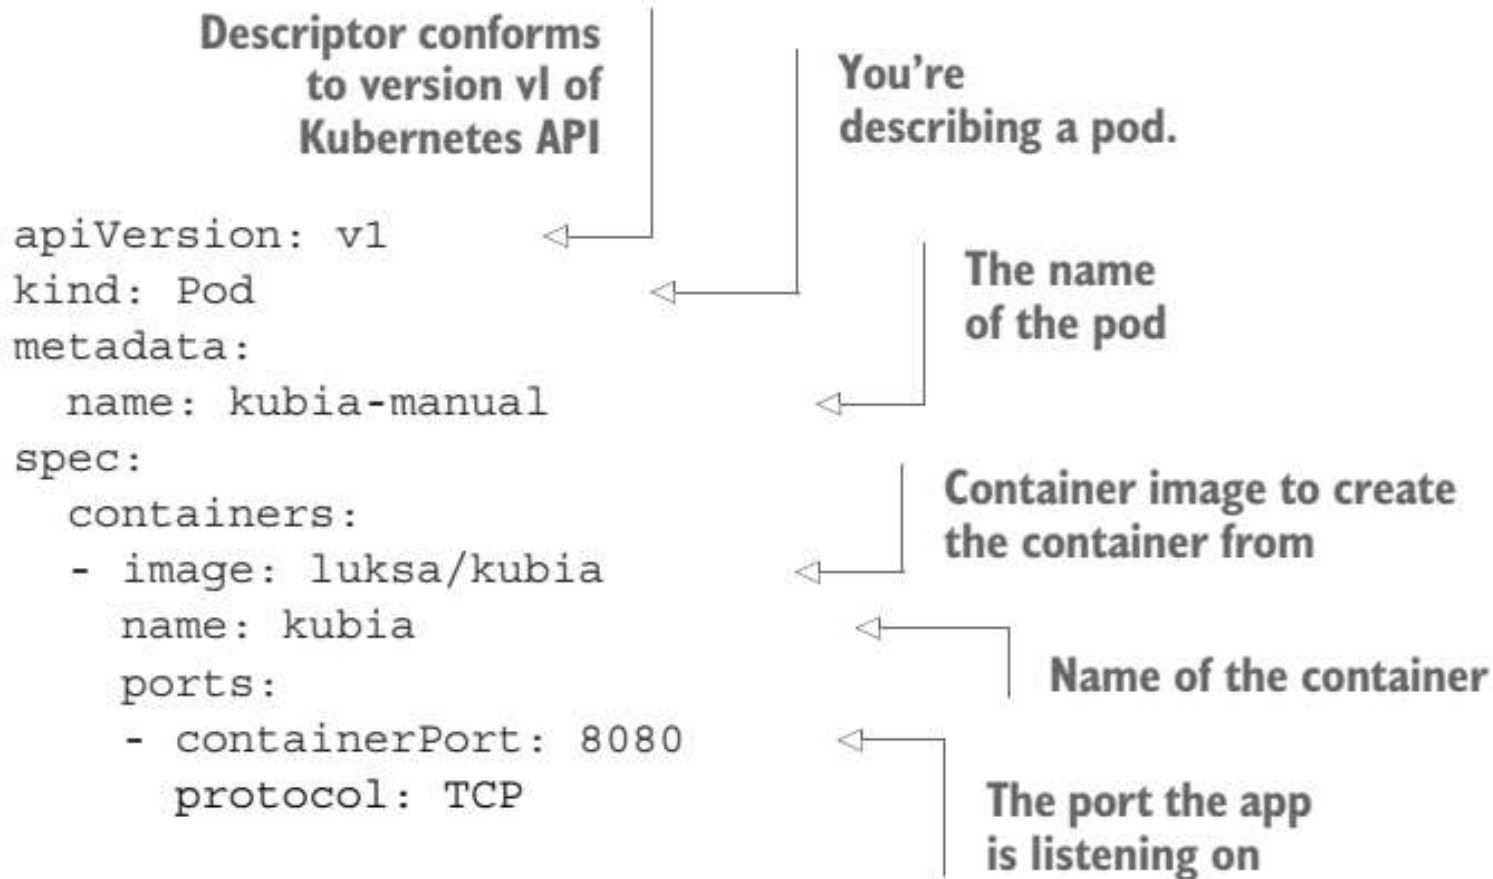



**Figure 2.8** Three instances of a pod managed by the same ReplicationController and exposed through a single service IP and port.



**Figure 3.2** Each pod gets a routable IP address and all other pods see the pod under that IP address.

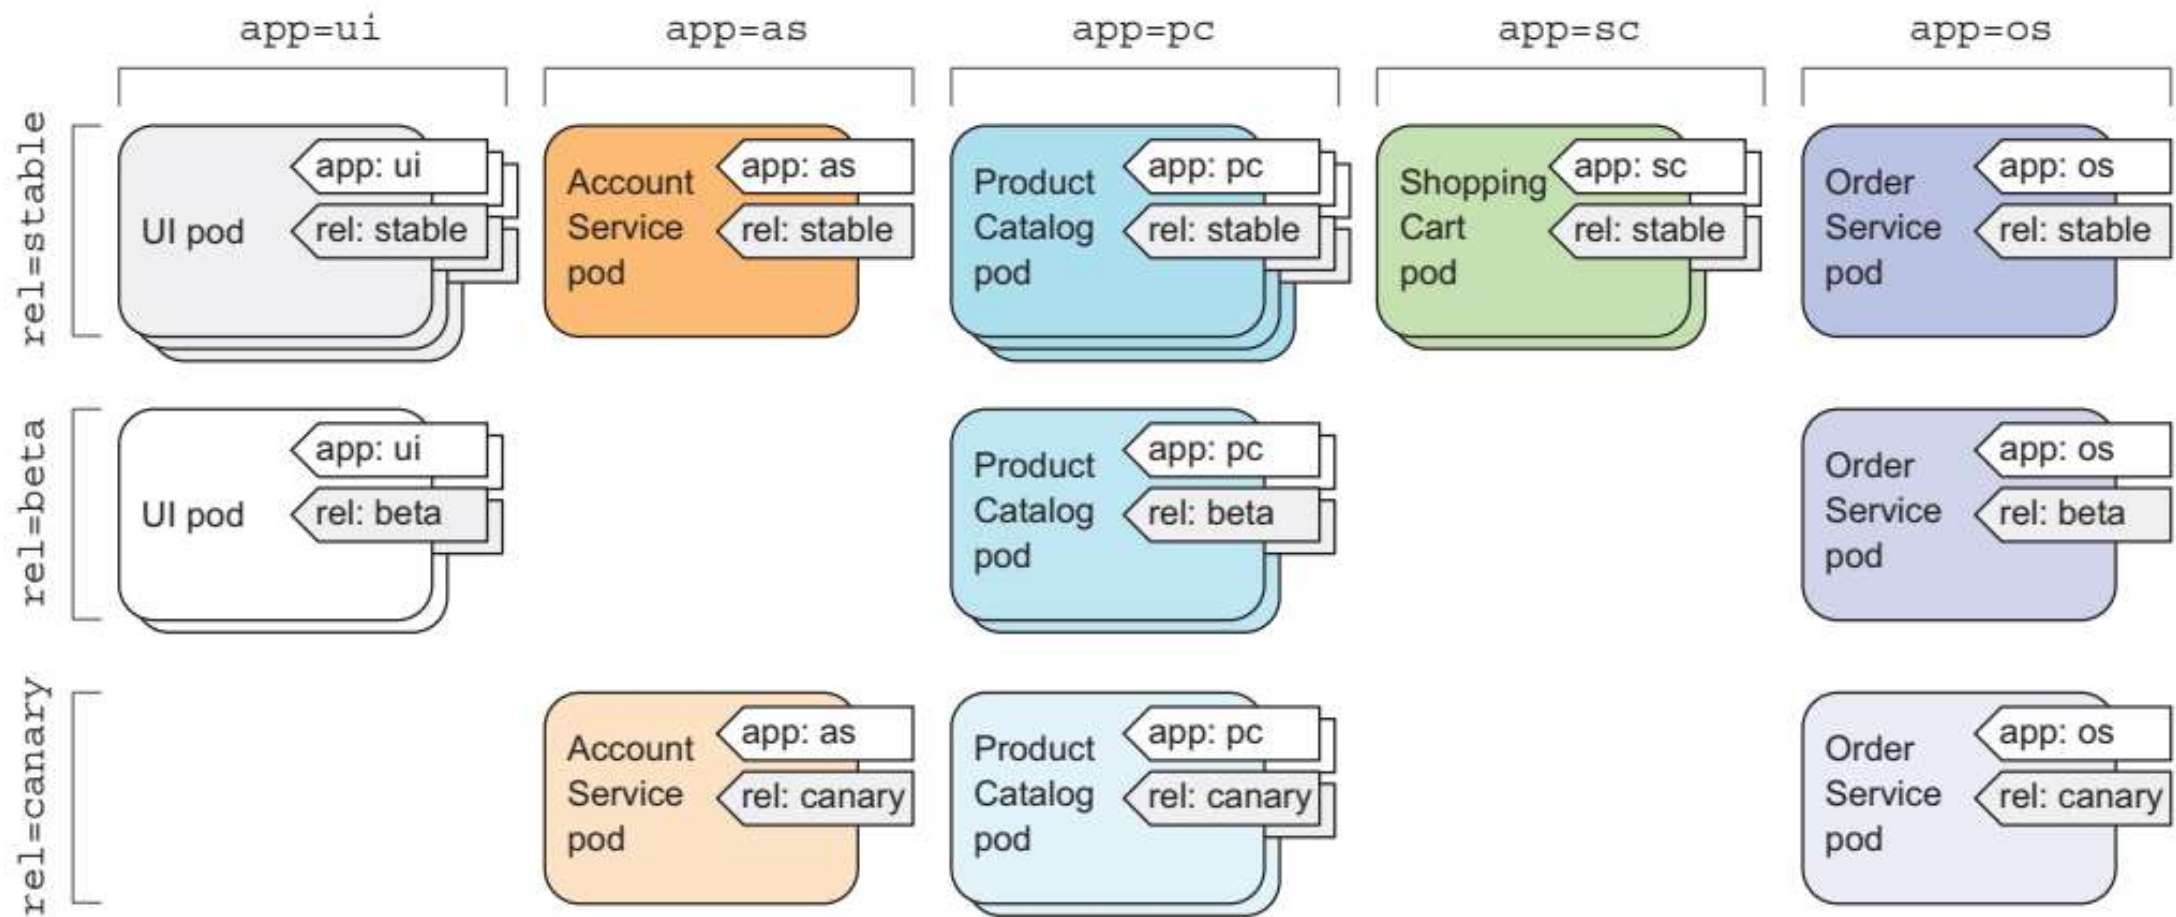
### Listing 3.2 A basic pod manifest: kuba-manual.yaml






**Figure 3.6** Uncategorized pods in a microservices architecture






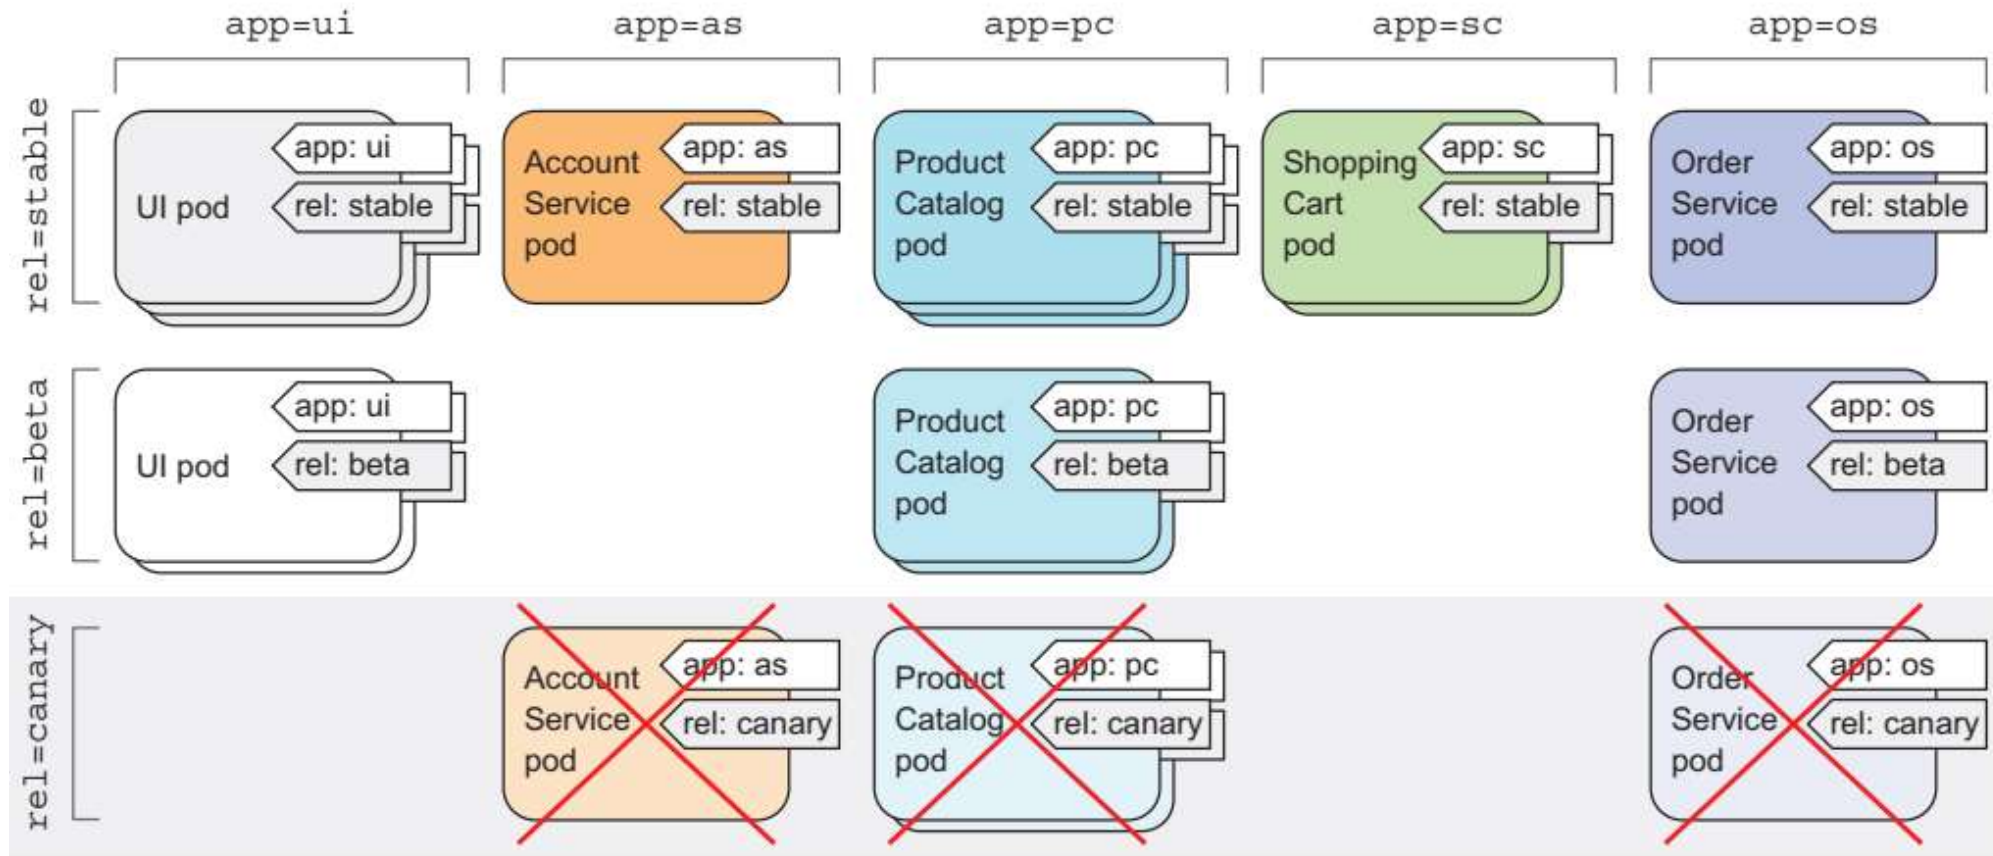
**Figure 3.7** Organizing pods in a microservices architecture with pod labels



**Figure 3.8** Selecting the product catalog microservice pods using the “app=pc” label selector




**Figure 3.9** Selecting pods with multiple label selectors




**Figure 3.10** Selecting and deleting all canary pods through the `rel=canary` label selector

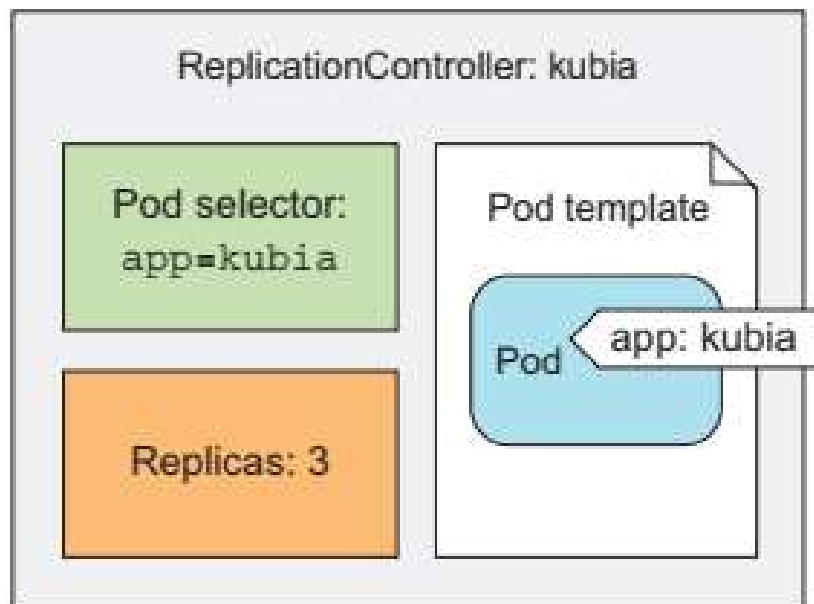




**Figure 4.1** When a node fails, only pods backed by a ReplicationController are recreated.



**Figure 4.2** A ReplicationController's reconciliation loop



**Figure 4.3** The three key parts of a `ReplicationController` (pod selector, replica count, and pod template)

#### Listing 4.4 A YAML definition of a ReplicationController: kubia-rc.yaml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: kubia
spec:
  replicas: 3
  selector:
    app: kubia
```

This manifest defines a  
ReplicationController (RC)

The name of this  
ReplicationController


The desired number  
of pod instances

The pod selector determining  
what pods the RC is operating on

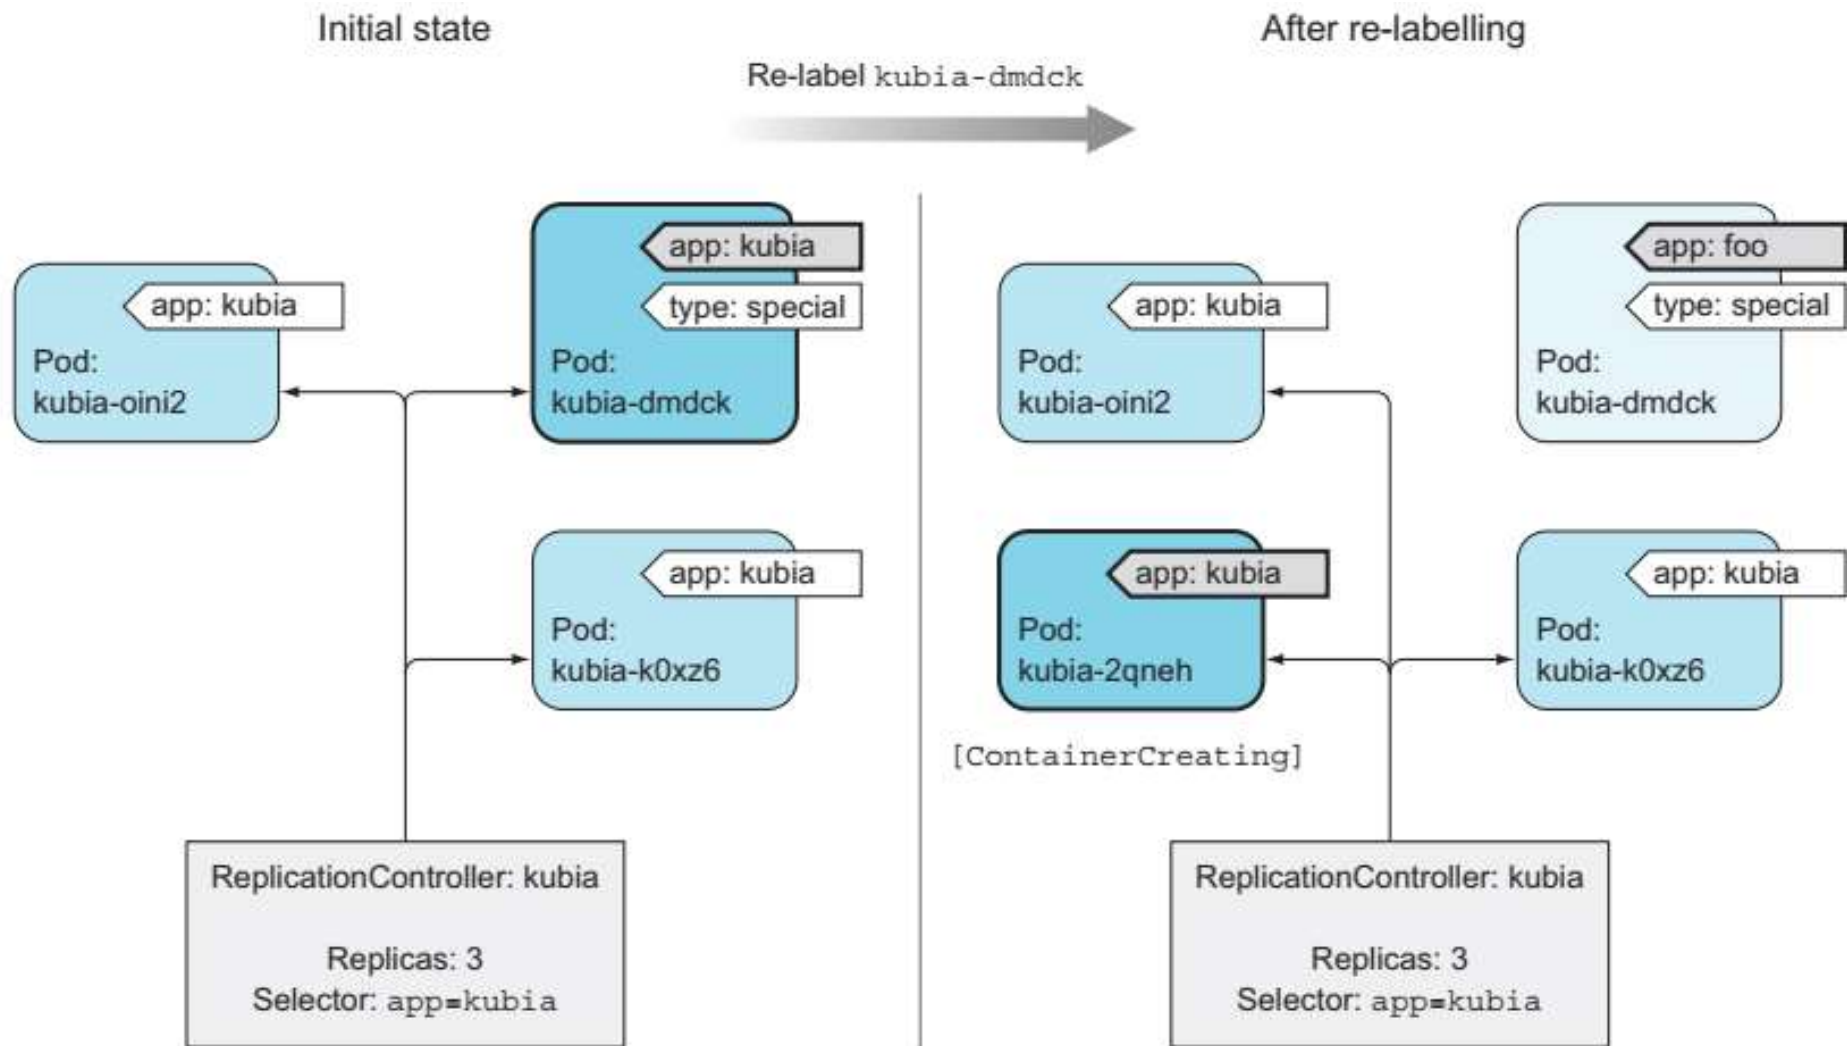
```
template:
  metadata:
    labels:
      app: kubia
  spec:
    containers:
      - name: kubia
        image: luksa/kubia
        ports:
          - containerPort: 8080
```

The pod template  
for creating new  
pods



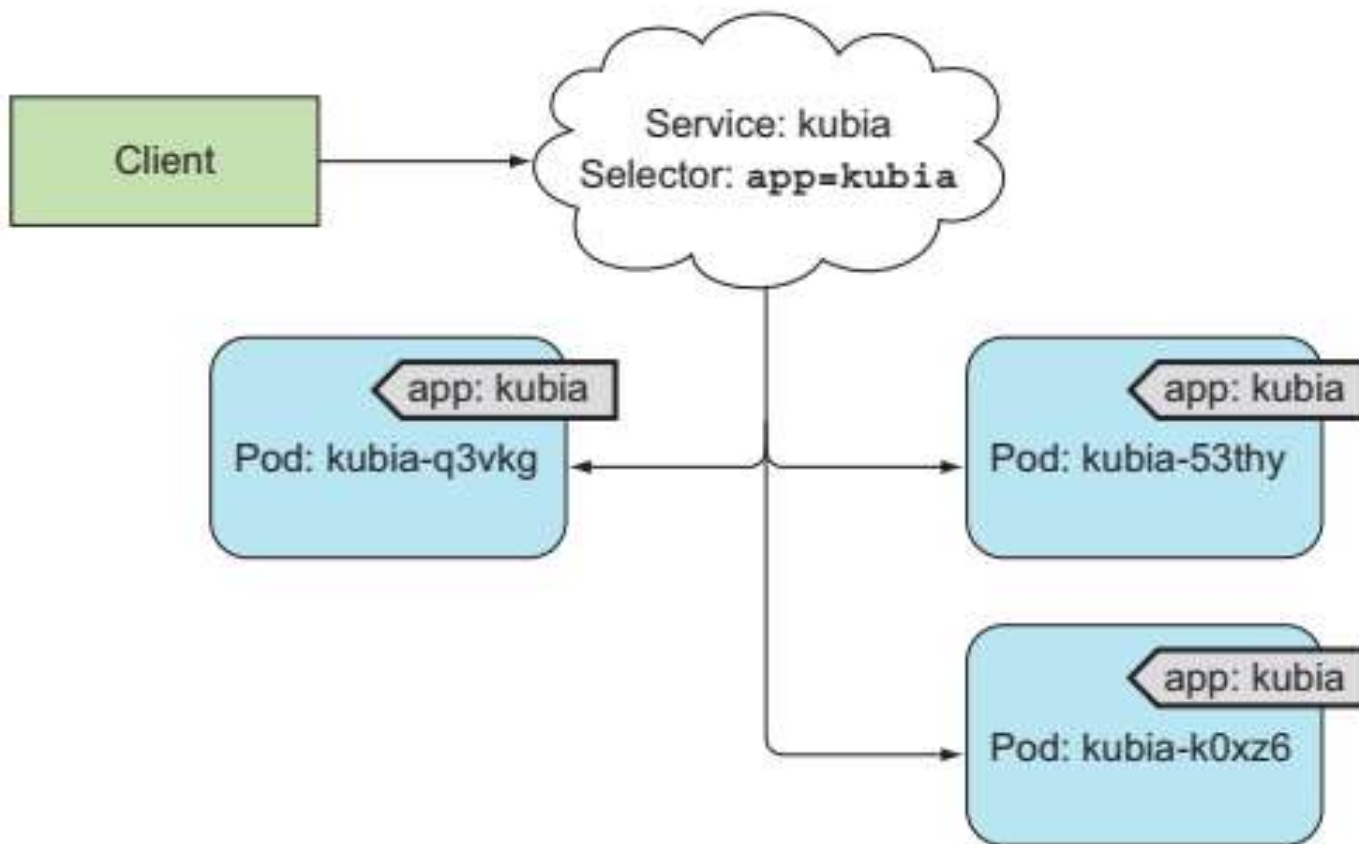


**Figure 4.4** If a pod disappears, the ReplicationController sees too few pods and creates a new replacement pod.



**Figure 4.5** Removing a pod from the scope of a ReplicationController by changing its labels

Services



**Figure 5.2** Label selectors determine which pods belong to the Service.


## Listing 5.1 A definition of a service: kubia-svc.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: kubia
spec:
  ports:
  - port: 80
    targetPort: 8080
  selector:
    app: kubia
```

The port this service  
will be available on


The container port the  
service will forward to

All pods with the app=kubia  
label will be part of this service.



**Figure 5.5** Exposing a service to external clients

# Kubernetes Architecture



# Kubernetes namespaces

- Provide for a scope of Kubernetes resource, carving up your cluster in smaller units
  - `$ kubectl get ns`
  - `$ kubectl describe ns default`
  - `$ kubectl create namespace test`



**Thanks**